

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

data = pd.read_csv('/content/events.csv')

data['event_time'] = pd.to_datetime(data['event_time'])

print(data.head())
print(data.info())
print(data.describe())

sns.countplot(x='event_type', data=data)
plt.title("Event Type Distribution")
plt.show()

top_brands = data['brand'].value_counts().head(10)
top_categories = data['category_code'].value_counts().head(10)

top_brands.plot(kind='bar', title="Top Brands", figsize=(10, 6))
plt.show()

top_categories.plot(kind='bar', title="Top Categories", figsize=(10, 6))
plt.show()

data['event_date'] = data['event_time'].dt.date
event_counts = data.groupby('event_date')['event_type'].count()
event_counts.plot(title="Events Over Time", figsize=(10, 6))
plt.show()

session_length = data.groupby('user_session')['event_time'].apply(lambda x: (x.max() - x.min()).seconds)
sns.histplot(session_length, bins=50, kde=True)
plt.title("Session Duration Distribution")
plt.xlabel("Session Duration (seconds)")
plt.show()
```



```
event_time event_type product_id category_id \
0 2020-09-24 11:57:06+00:00 view 1996170 2144415922528452715
1 2020-09-24 11:57:26+00:00 view 139905 2144415926932472027
2 2020-09-24 11:57:27+00:00 view 215454 2144415927158964449
3 2020-09-24 11:57:33+00:00 view 635807 2144415923107266682
4 2020-09-24 11:57:36+00:00 view 3658723 2144415921169498184
```

```
category_code brand price user_id \
0 electronics.telephone NaN 31.90 1515915625519388267
1 computers.components.cooler zalman 17.16 1515915625519380411
2 NaN NaN 9.81 1515915625513238515
3 computers.peripherals.printer pantum 113.81 1515915625519014356
4 NaN caméronsino 15.87 1515915625510743344
```

```
user_session
0 LJJuVLEjPT
1 tdicluNnRY
2 4TMArHtXQy
3 aGFYrNgC08
4 aa4mmk0kwQ
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 885129 entries, 0 to 885128
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null	Count	Dtype
0	event_time	885129 non-null		datetime64[ns, UTC]
1	event_type	885129 non-null		object
2	product_id	885129 non-null		int64
3	category_id	885129 non-null		int64
4	category_code	648910 non-null		object
5	brand	672765 non-null		object
6	price	885129 non-null		float64
7	user_id	885129 non-null		int64
8	user_session	884964 non-null		object

```
dtypes: datetime64[ns, UTC](1), float64(1), int64(3), object(4)
```

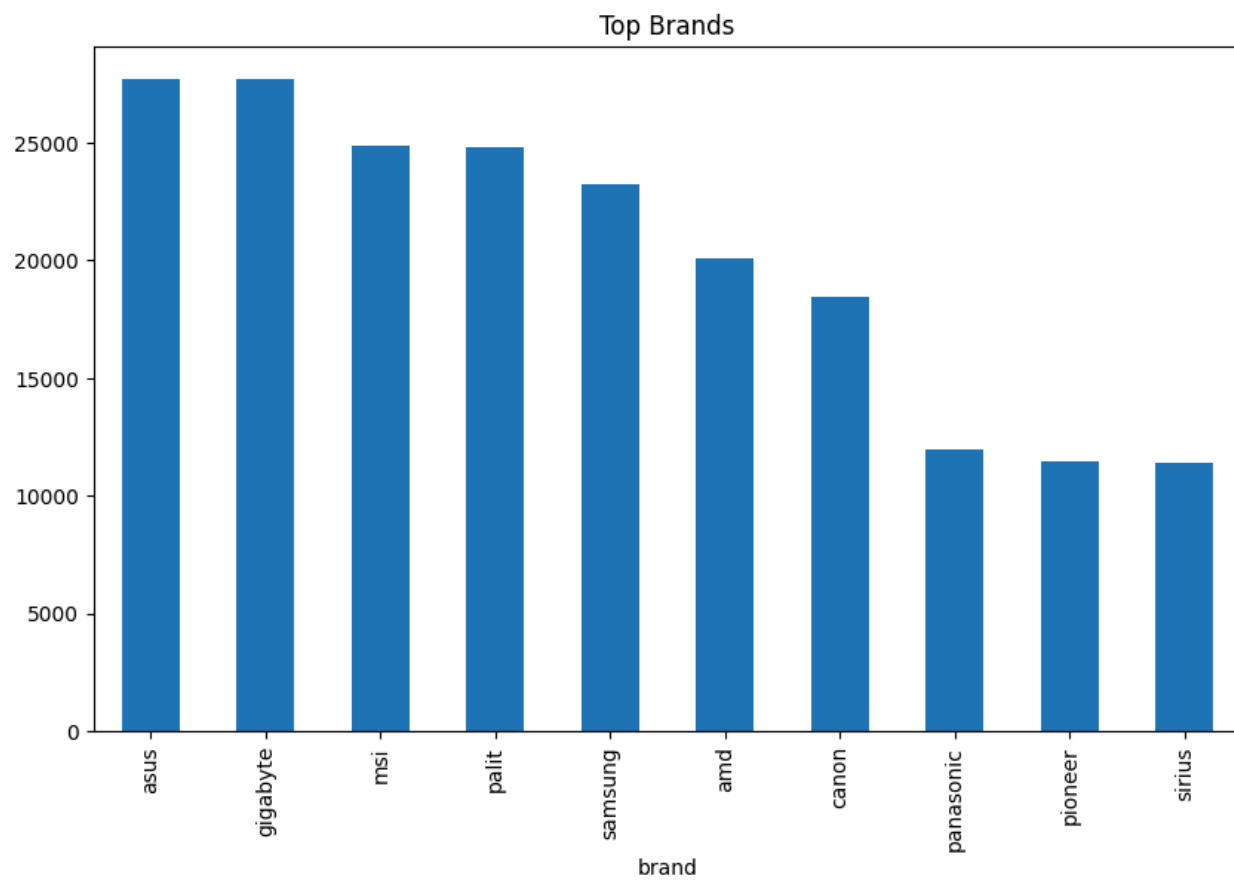
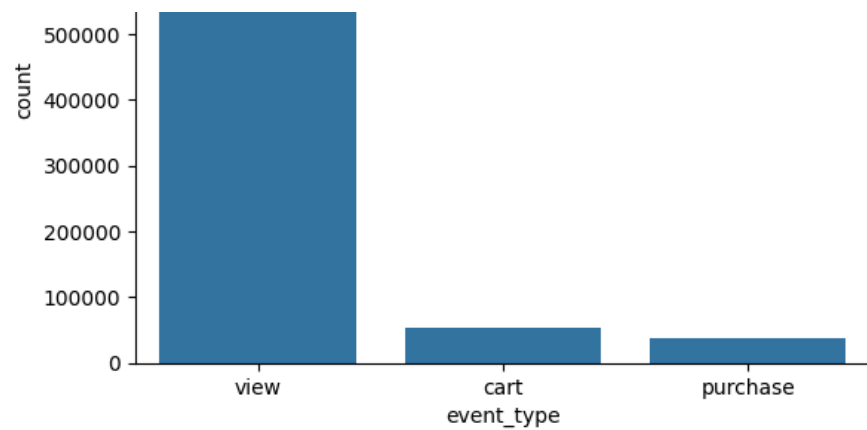
```
memory usage: 60.8+ MB
```



```
None
```


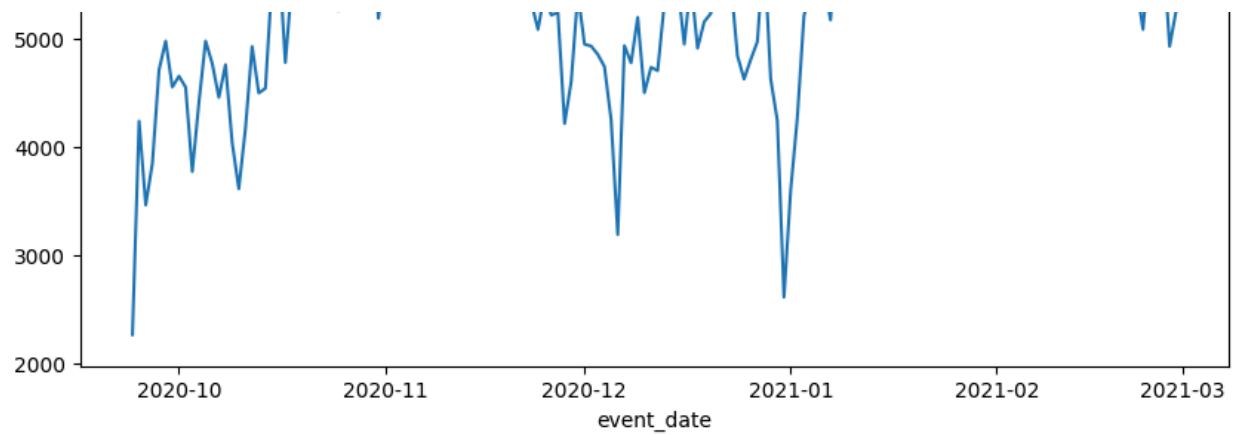
	product_id	category_id	price	user_id
count	8.851290e+05	8.851290e+05	885129.000000	8.851290e+05
mean	1.906621e+06	2.144423e+18	146.328713	1.515916e+18
std	1.458708e+06	6.165105e+14	296.807683	3.747287e+07
min	1.020000e+02	2.144416e+18	0.220000	1.515916e+18
25%	6.988030e+05	2.144416e+18	26.460000	1.515916e+18
50%	1.452883e+06	2.144416e+18	65.710000	1.515916e+18
75%	3.721194e+06	2.144416e+18	190.490000	1.515916e+18
max	4.183880e+06	2.227847e+18	64771.060000	1.515916e+18

Event Type Distribution









```

from datetime import timedelta

# Define churn: No purchase in the last 30 days
data['is_purchase'] = data['event_type'] == 'purchase'
last_purchase_date = data.loc[data['is_purchase']].groupby('user_id')['event_time'].max()
max_date = data['event_time'].max()

# Churn threshold: 30 days since last purchase
churn_threshold = last_purchase_date + timedelta(days=30)
data['churn_status'] = data['user_id'].map(lambda x: x not in churn_threshold.index or churn_threshold[x] < max_date)

# Verify churn rate
churn_rate = data['churn_status'].mean()
print("The churn rate is {:.2%}".format(churn_rate))

```

➡ The churn rate is 94.90%

```

# RFM Metrics
data['recency'] = data.groupby('user_id')['event_time'].transform(lambda x: (max_date - x.max()).days)
data['frequency'] = data.groupby('user_id')['event_type'].transform('count')
data['monetary'] = data.groupby('user_id')['price'].transform('sum')

# Behavioral patterns
data['view_to_cart_ratio'] = data.groupby('user_id')['event_type'].apply(
    lambda x: (x == 'cart').sum() / max((x == 'view').sum(), 1)
)
data['cart_to_purchase_ratio'] = data.groupby('user_id')['event_type'].apply(
    lambda x: (x == 'purchase').sum() / max((x == 'cart').sum(), 1)
)

# Preferences
data['top_brand'] = data.groupby('user_id')['brand'].transform(lambda x: x.mode()[0] if not x.mode().empty else 'Unknown')
data['top_category'] = data.groupby('user_id')['category_code'].transform(lambda x: x.mode()[0] if not x.mode().empty else 'Unknown')

# Final user-level dataset
user_features = data.groupby('user_id').agg({
    'recency': 'first',
    'frequency': 'first',
    'monetary': 'first',
    'view_to_cart_ratio': 'first',
    'cart_to_purchase_ratio': 'first',
    'top_brand': 'first',
    'top_category': 'first',
    'churn_status': 'first'
}).reset_index()

print(user_features.head())

```

```

user_id  recency  frequency  monetary  view_to_cart_ratio \
0  1515915625353226922    122         1      76.48         NaN
1  1515915625353230067    145         1      28.98         NaN
2  1515915625353230683     78        13     814.93         NaN
3  1515915625353230922    149         1     274.40         NaN
4  1515915625353234047     10        36    5481.90         NaN

```

```

cart_to_purchase_ratio  top_brand  top_category \
0         NaN    honor    electronics.clocks
1         NaN    kester    Unknown
2         NaN    creative    electronics.audio.acoustic
3         NaN    msi    computers.components.videocards
4         NaN    samsung    electronics.audio.headphone

```

```

churn_status
0         True
1         True
2         True
3         True
4         True

```

```

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score, log_loss, brier_score_loss

# Prepare data for modeling
X = user_features[['recency', 'frequency', 'monetary', 'view_to_cart_ratio', 'cart_to_purchase_ratio', 'top_brand', 'top_category']]
X = pd.get_dummies(X, columns=['top_brand', 'top_category'], drop_first=True)
y = user_features['churn_status']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Random Forest Classifier
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Model Evaluation
y_pred = model.predict(X_test)
y_prob = model.predict_proba(X_test)[:, 1]

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# AUC Score
auc_score = roc_auc_score(y_test, y_prob)
print(f"AUC Score: {auc_score:.4f}")

# Log Loss
log_loss_value = log_loss(y_test, model.predict_proba(X_test))
print(f"Log Loss: {log_loss_value:.4f}")

```

```
# Brier Score Loss
brier_loss_value = brier_score_loss(y_test, y_prob)
print(f"Brier Score Loss: {brier_loss_value:.4f}")
import shap

# Explain model predictions using SHAP
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_test)

# SHAP Summary Plot
shap.summary_plot(shap_values[1], X_test)
```

🔄 Classification Report:

	precision	recall	f1-score	support
False	0.43	0.29	0.35	886
True	0.99	1.00	0.99	80571
accuracy			0.99	81457
macro avg	0.71	0.64	0.67	81457
weighted avg	0.99	0.99	0.99	81457

AUC Score: 0.9670